

Intro

Please implement **one** of the following exercises, share your code (on GitHub) with your Pulumi recruiter prior to your HW Review, and be prepared to walk through it on a call. If any details of the exercise are unclear, please ask for clarification.

Please timebox your work on the exercise to **four hours or less**. It is acceptable to return to an incomplete/non-working exercise.

Option 1 - Container Application (non-Kubernetes)

Using Pulumi, create and deploy a web application running in a managed container service such as AWS ECS, Azure Container Apps, or GCP Cloud Run. The web application should display a customized web page that returns a configurable value in its web page response.

Use **one** of AWS, Azure, or Google Cloud and [your language of choice](#).

At a minimum your Pulumi application code should consist of:

- A Dockerfile to build your custom docker image
 - Suggestion: use [this gist](#) as your starting app and Docker image. Modify the application to read a configurable value via an environment variable.
- A managed container service, including but not limited to:
 - AWS ECS
 - Azure Container Apps
 - Azure Container Instances
 - GCP Cloud Run
- Relevant cloud provider resources to deploy and run the Docker image.
- A configurable value (via [pulumi config](#)) that changes the value displayed in the web page response.

NOTE: you may not utilize any ECS functionality provided by the [Pulumi AWSX](#) package

Success criteria:

- Should be able to deploy the provided Web Application into your network.

Option 2 - Kubernetes Application

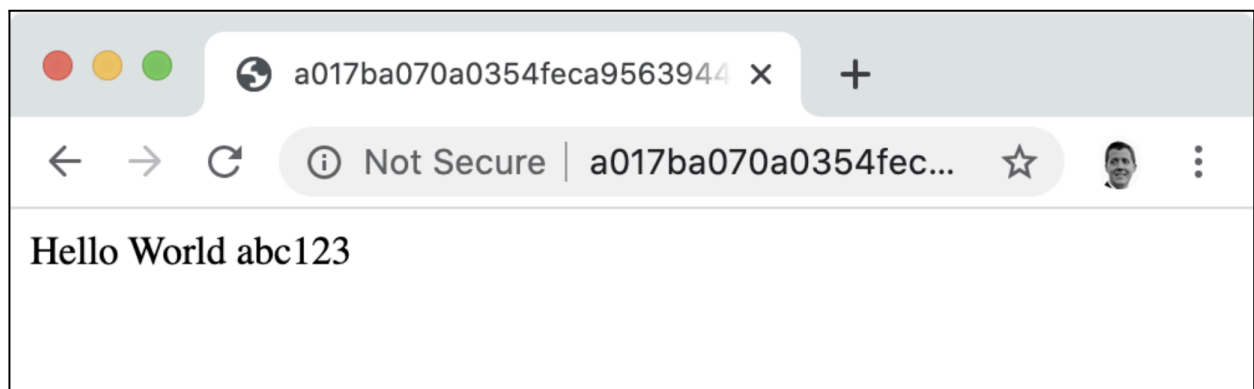
Using Pulumi, create and deploy a web application running in Kubernetes. The web application should display a customized web page that returns a configurable value in its web page response.

Use **one** of AWS, Azure, Google Cloud, or Digital Ocean and [your language of choice](#).

At a minimum your Pulumi application code should consist of:

- A Dockerfile to build your custom docker image
 - Suggestion: use [this gist](#) as your starting app and Docker image. Modify the application to read a configurable value via an environment variable.
- A Kubernetes cluster
- Relevant Kubernetes resources to deploy and run the Docker image
- A configurable value (via [pulumi config](#)) that changes the value displayed in the web page response

For clarity, deploy a web application that serves a web page like below where abc123 is the configurable value:



Success criteria:

- Should be able to load the web page in a browser
- Should be able to change the configured value, run `pulumi up`, reload the web page and see the value change

Option 3 - Web Server Fleet Abstraction

Using Pulumi, create an abstraction for a “web server fleet” that supports at least two operating systems. Your application should accept the following (pseudocode) as input and provision virtual machines as defined by the inputs.

```
new WebServerFleet({
  subnets: ['subnet-abc123', 'subnet-abc123'],
  machines: [
    {os: 'ubuntu', size: 'small', count: 3},
    {os: 'amazonlinux', size: 'medium', count: 2},
  ],
})
```

Use **one** of AWS, Azure, Google Cloud, or Digital Ocean and [your language of choice](#).

At a minimum your Pulumi application code should consist of:

- A minimal set of resources to build a VPC/network on the cloud you chose (no abstraction needed)
- An interface/class for the arguments for your web server fleet that takes the following:
 - The network and/or subnets to deploy into
 - The number of virtual machines to deploy
 - The “abstract” size of the virtual machines in the form of:
 - “small”
 - “medium”
 - “large”
 - The “abstract” operating system of the virtual machines in the form of:
 - “ubuntu”
 - “amazon”
 - (or other OS of your choice)
- A `WebServerFleet` abstraction via a [Component Resource](#) implementation that does the following:
 - Provisions the defined number of VMs meeting using the provided inputs
 - Picks a different cloud machine type based on the abstract virtual machine size, e.g.:
 - “small” => “t2.medium”
 - “medium” => “t2.large”
 - “large” => “t2.2xlarge”
- Picks the appropriate OS based on the OS type input, e.g.
 - “ubuntu” => “ubuntu-18.04”
 - “amazon” => “amzn2-ami-hvm-2.0”

- For each operating system implementation install a web server using a startup script via userdata or metadata startup script for the virtual machine.
 - Suggestion: use nginx for simplicity.

Success criteria:

- A minimal set of resources to build a VPC/network
- Given the WebServerFleet inputs, deploys the correct type and number of virtual machines specified
- Note: Ubuntu and Amazon Linux are just examples, use the operating systems of your choice

